BIG DATA WEEK 11 TASK REPORT

(JOBSHEET MAPREDUCE 3)

By:

Rajendra Rakha Arya Prabaswara

1941720080

(TI-3H / 20)

INFORMATICS ENGINEERING STUDY PROGRAM MAJORING INFORMATION TECHNOLOGY

STATE POLYTECHNIC MALANG

Part I : Creating a MapReduce Job with JetBrains IntelliJ IDEA

1. Install and will be able to access the Hadoop cluster correctly .

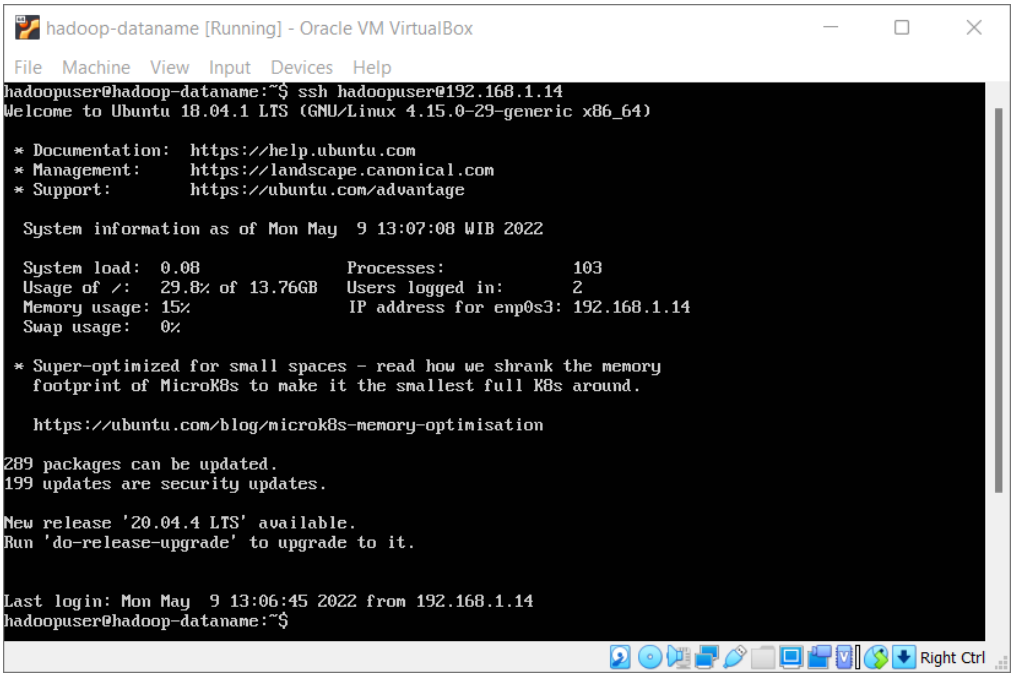

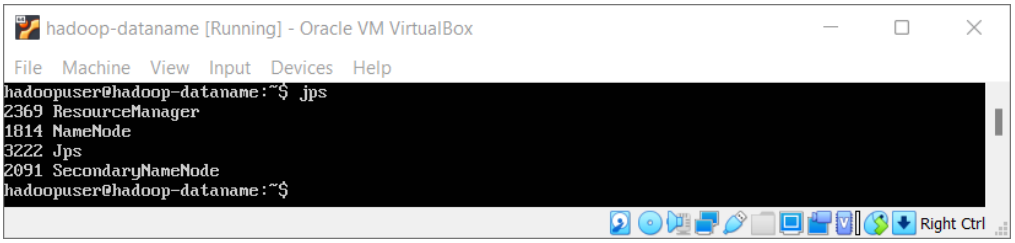

2. Open the IntelliJ IDEA IDE IDE and create a new project with the Java Language and maven framework build. And select the archetype "maven quick-start".

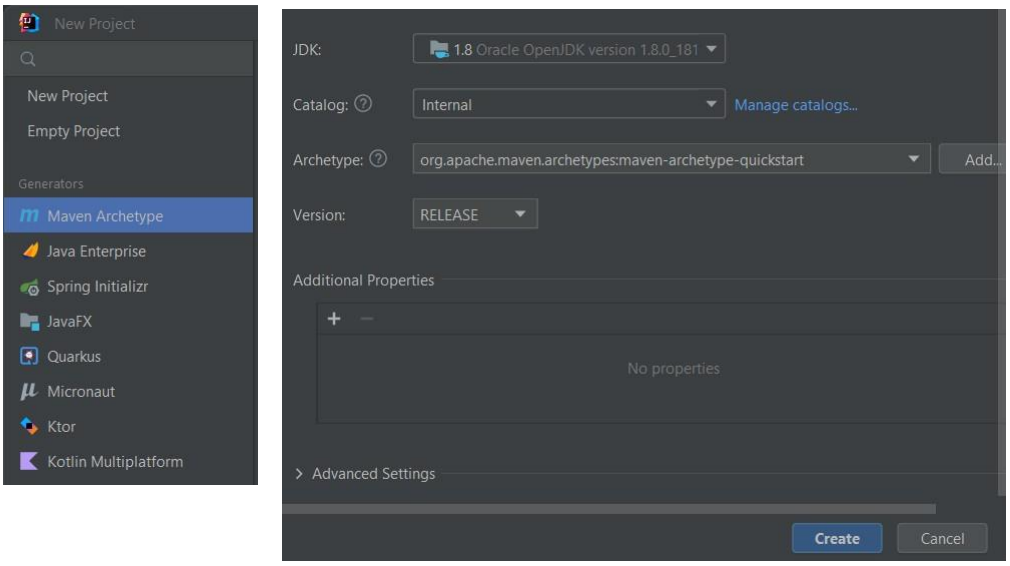

3. After the project is completed. In pom.xml added dependencies in order to create a MapReduce Job. Reload the project by clicking on the 'm' icon.

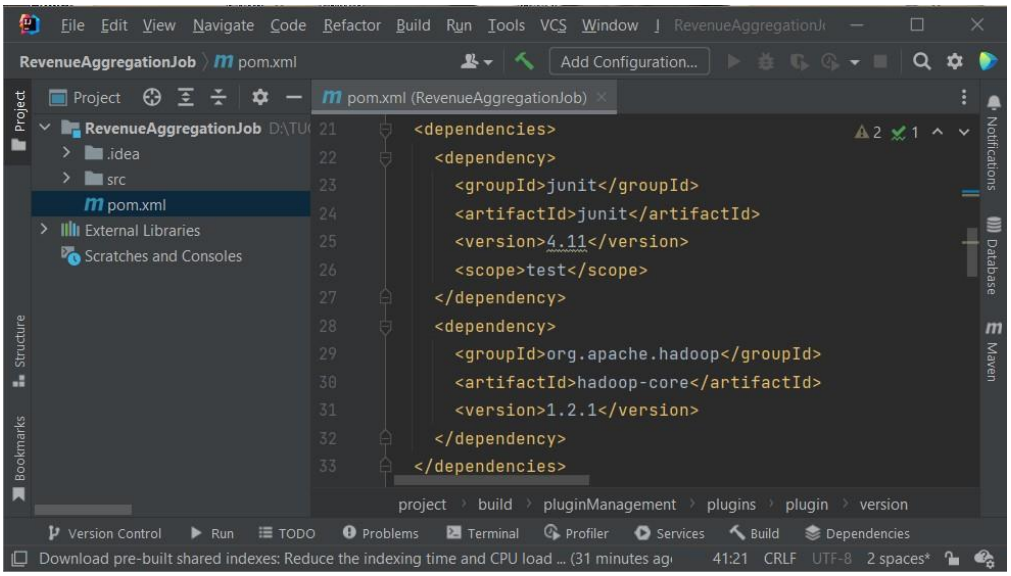

4. Create 2 new classes on the project under the names RevenueAggregationMapper and RevenueAggregationReducer.

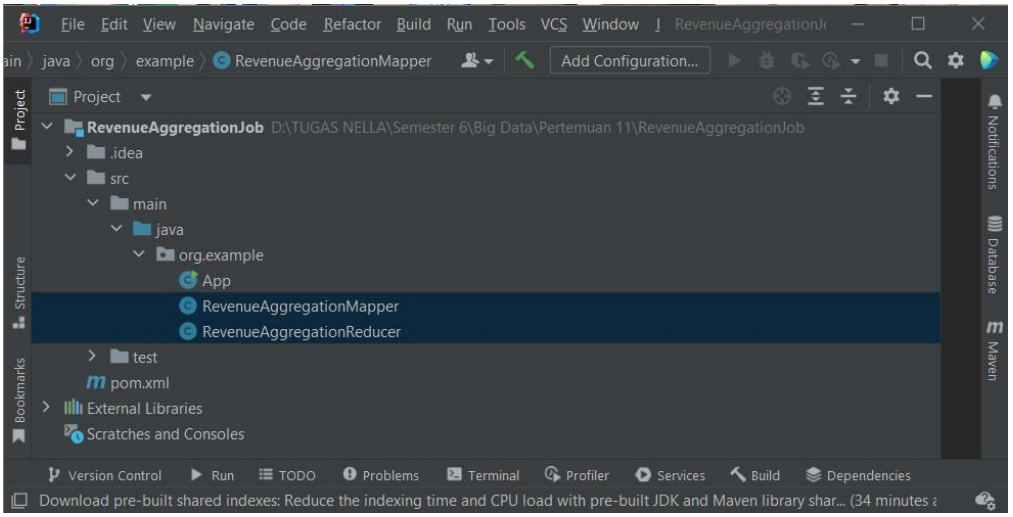

5. Modify the Class RevenueAggregationMapper.

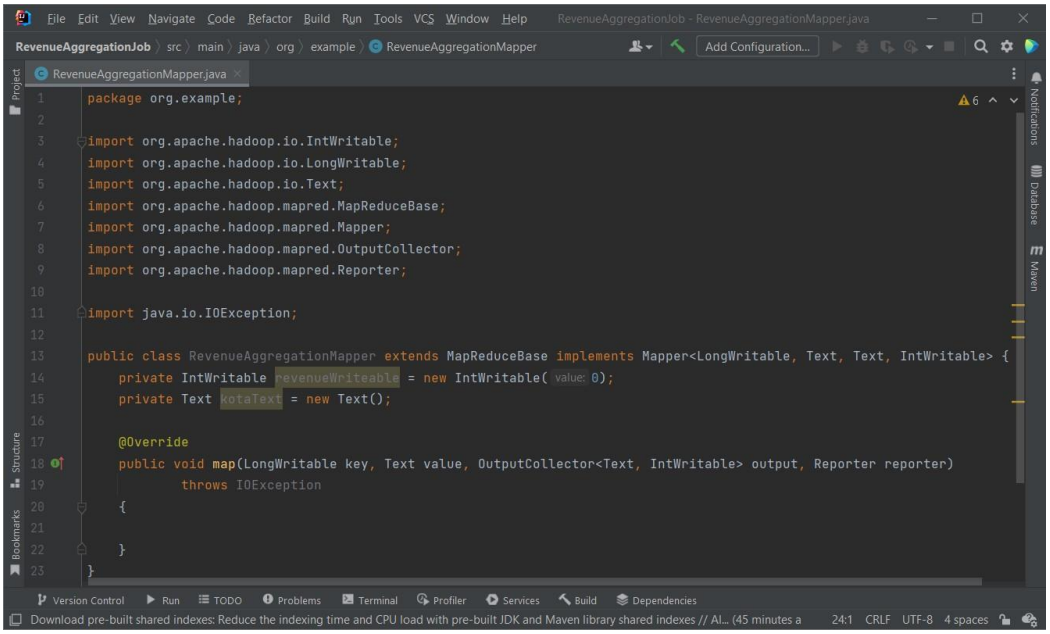

6. Modify the Class RevenueAggregationReducer.

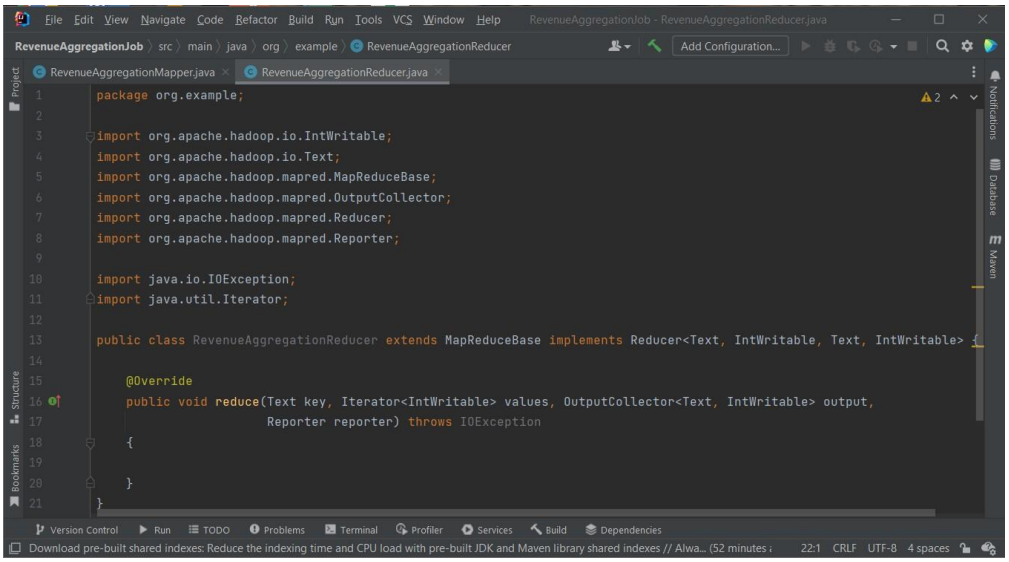

```
package org.example;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;

import java.io.IOException;
import java.util.Iterator;

public class RevenueAggregationReducer extends MapReduceBase implements Reducer<Text, IntWritable, Text, IntWritable> {

    @Override
    public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable> output,
                       Reporter reporter) throws IOException
    {

    }
}
```

7. Add imports blocks to the App class that is the project's built-in class.

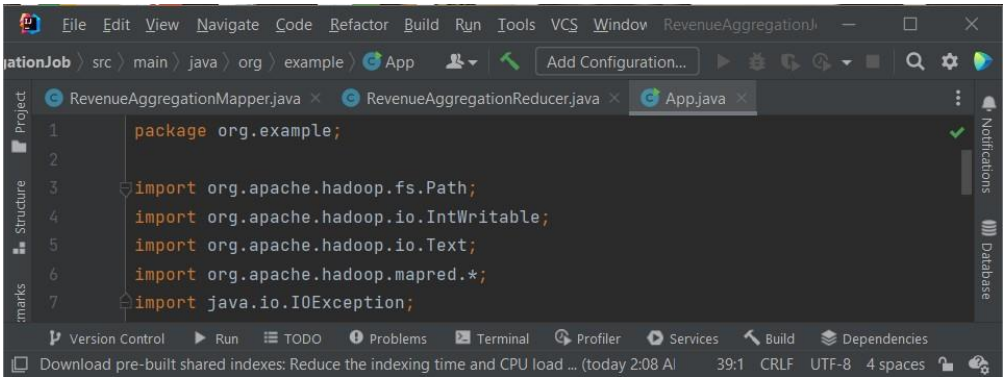

```
package org.example;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;
import java.io.IOException;
```

8. Then, add the code to the same class body .

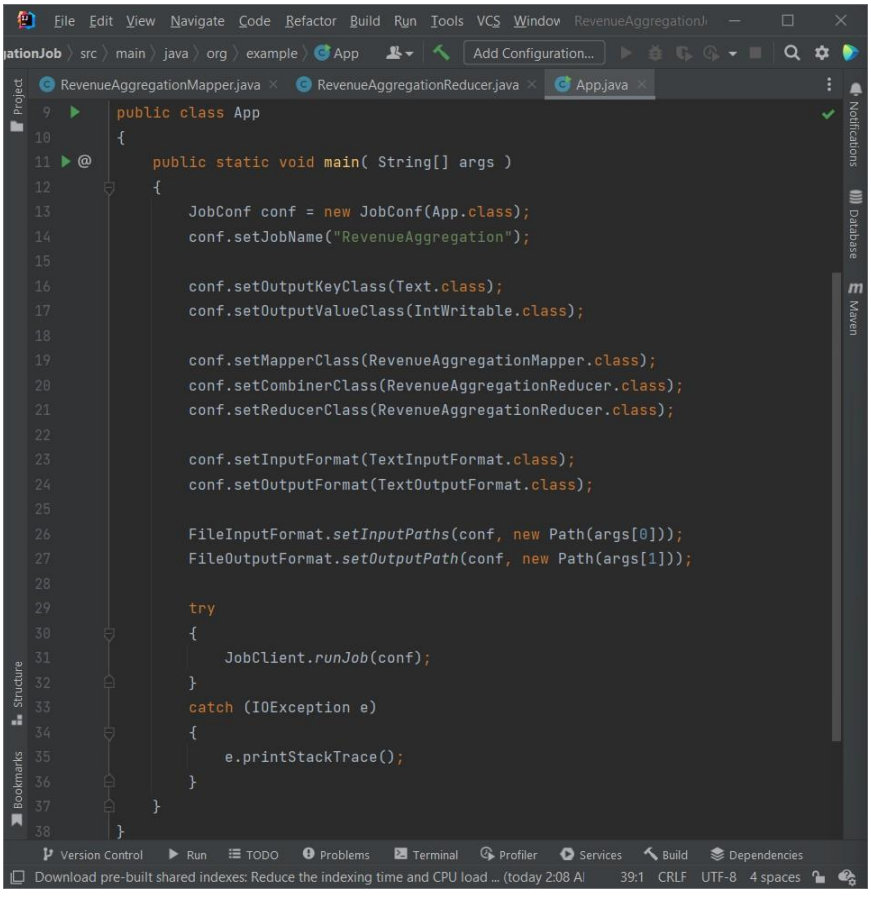

```
public class App
{
    public static void main( String[] args )
    {
        JobConf conf = new JobConf(App.class);
        conf.setJobName("RevenueAggregation");

        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);

        conf.setMapperClass(RevenueAggregationMapper.class);
        conf.setCombinerClass(RevenueAggregationReducer.class);
        conf.setReducerClass(RevenueAggregationReducer.class);

        conf.setInputFormat(TextInputFormat.class);
        conf.setOutputFormat(TextOutputFormat.class);

        FileInputFormat.setInputPaths(conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(conf, new Path(args[1]));

        try
        {
            JobClient.runJob(conf);
        }
        catch (IOException e)
        {
            e.printStackTrace();
        }
    }
}
```

9. Compile the program using the help of Maven by opening the Maven panel located to the right of the IDE. In the pane expand the folder "Lifecycle", select "Clean" and "Installl". Then click Run.

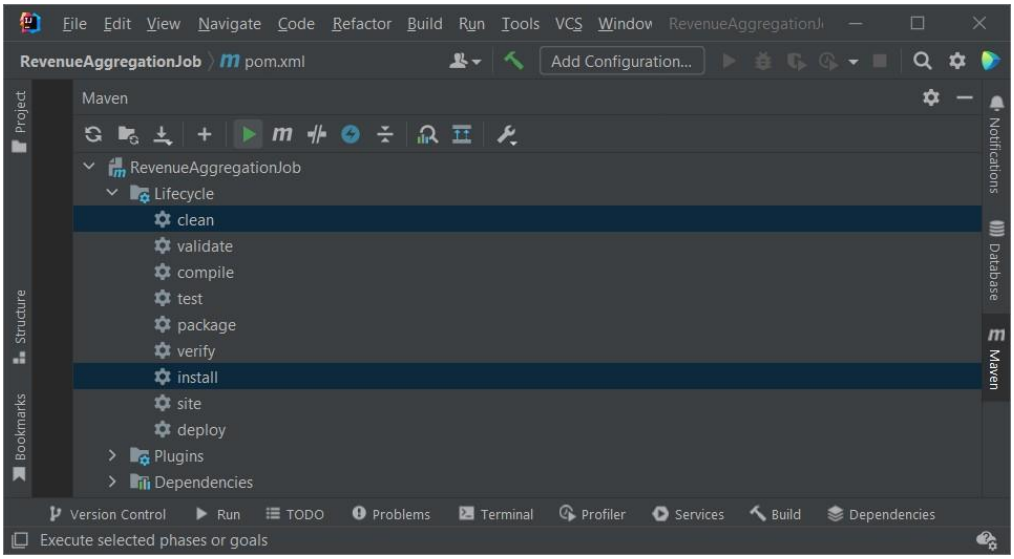

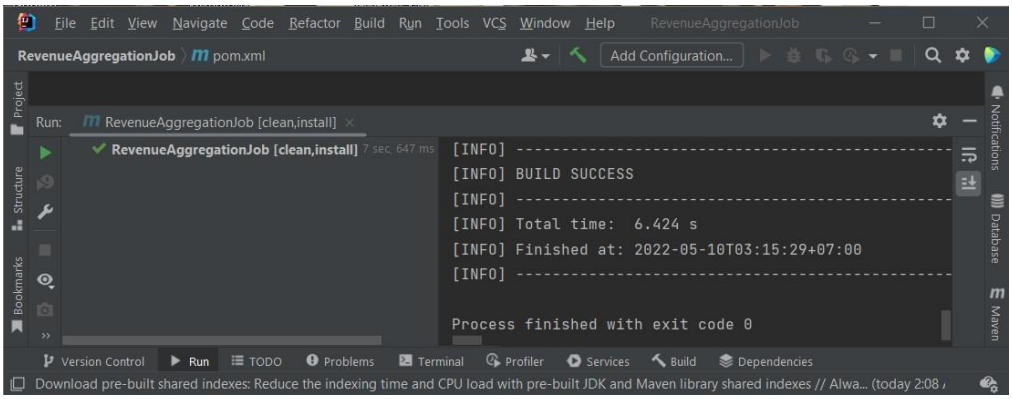

10. Create a plain text file named input01.

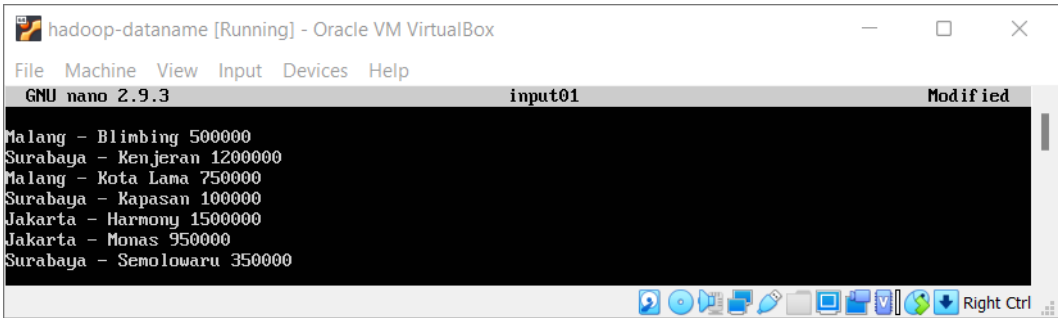

Upload the file to HDFS in the /YourName/Input folder. Also create 1 other folder in HDFS under the name /NamAnda/Output.

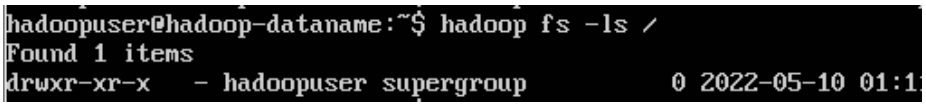

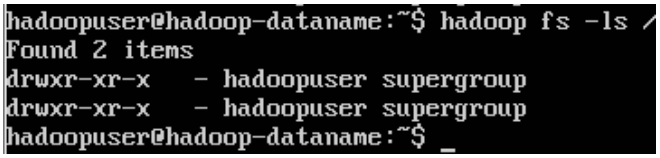

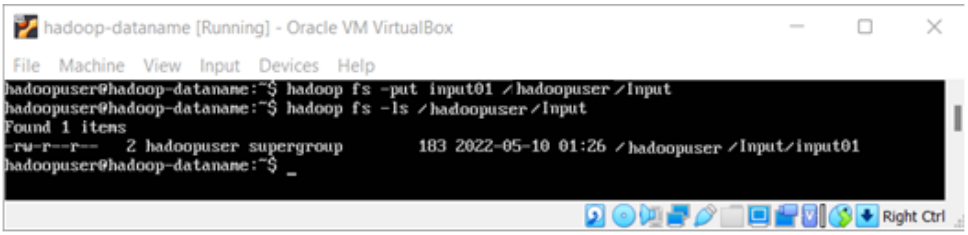

11. Upload jar files to the Name Node by using SCP or other methods. Execute the JAR using the command: Hadoop jar <nama_file>.jar <package_identifier> <folder_input_di_hdfs> <folder_output_di_hdfs>. JAR files uploaded to the namenode aredibe rinama revenue.jar.

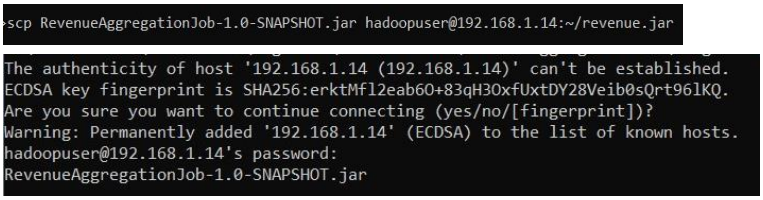

12. When the execution process is complete, the folder /Your Name/Output/output/result in which there is a part-00000 file that when you

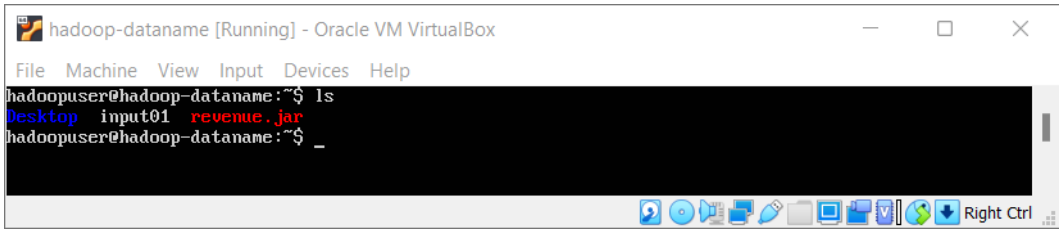

read, the contents are still empty. [Question] Why is the output file empty/has no content?

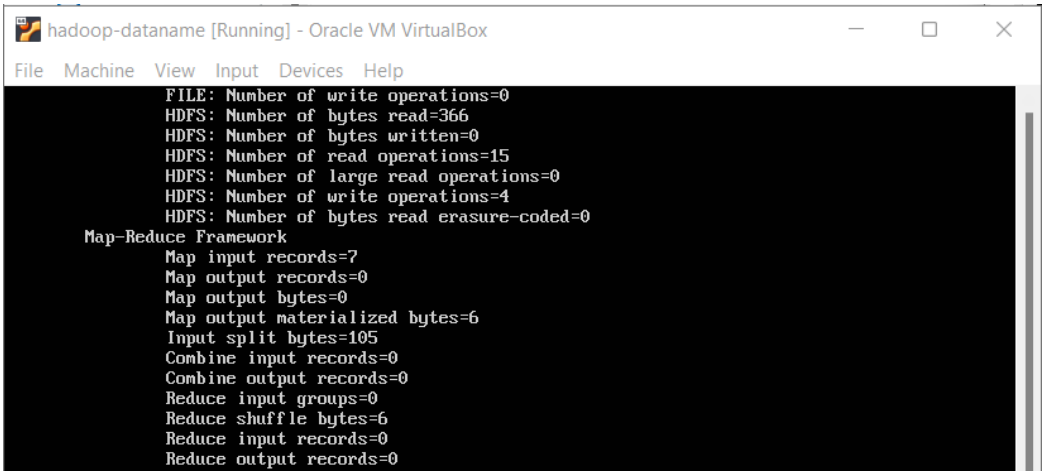

```
                Spilled Records=0
                Shuffled Maps =1
                Failed Shuffles=0
                Merged Map outputs=1
                GC time elapsed (ms)=22
                Total committed heap usage (bytes)=335683584
        Shuffle Errors
                BAD_ID=0
                CONNECTION=0
                IO_ERROR=0
                WRONG_LENGTH=0
                WRONG_MAP=0
                WRONG_REDUCE=0
        File Input Format Counters
                Bytes Read=183
        File Output Format Counters
                Bytes Written=0
hadoopuser@hadoop-dataname:~$
```

Right Ctrl

hadoop-dataname [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

```
hadoopuser@hadoop-dataname:~$ hadoop fs -ls
Found 1 items
drwxr-xr-x   - hadoopuser supergroup          0 2022-05-10 04:02
hadoopuser@hadoop-dataname:~$
```

Right Ctrl

hadoop-dataname [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

```
hadoopuser@hadoop-dataname:~$ hadoop fs -cat /Output/Hasil/part-00000
hadoopuser@hadoop-dataname:~$ _
```

Right Ctrl

Bagian II : Mapper & Reducer

13. Modify the RevenueAggregationMapper class to find out how many times Mapper is run.

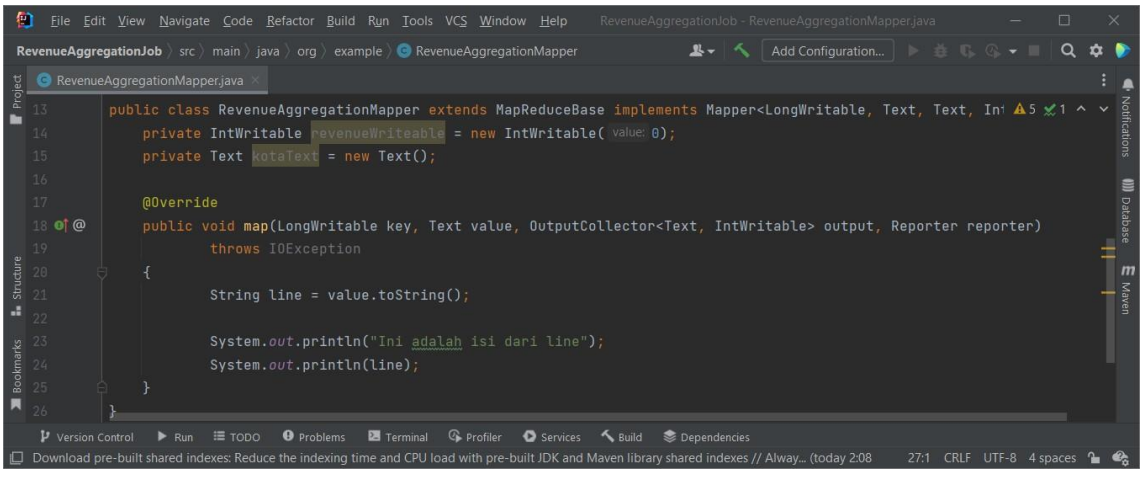

14. Modify the Class RevenueAggregationReducer.

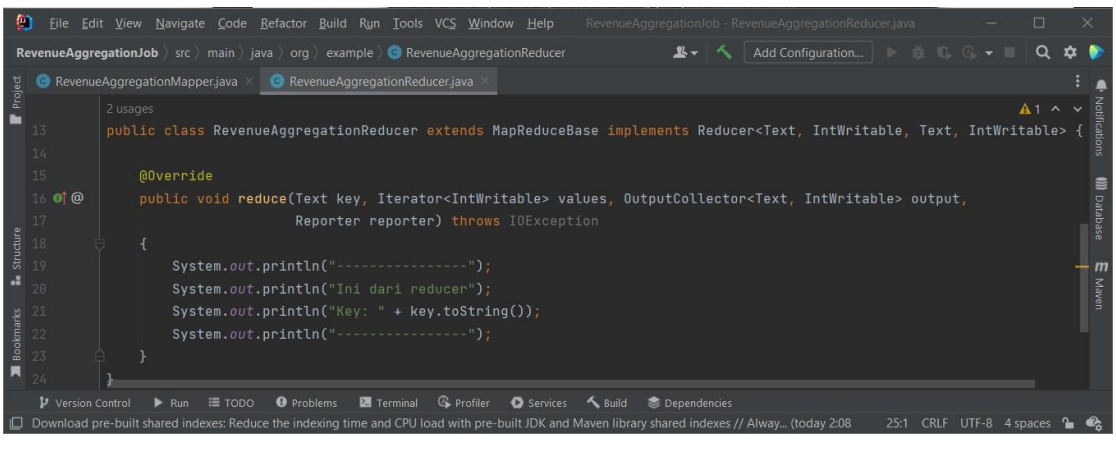

15. Rebuild jar again by using Maven.

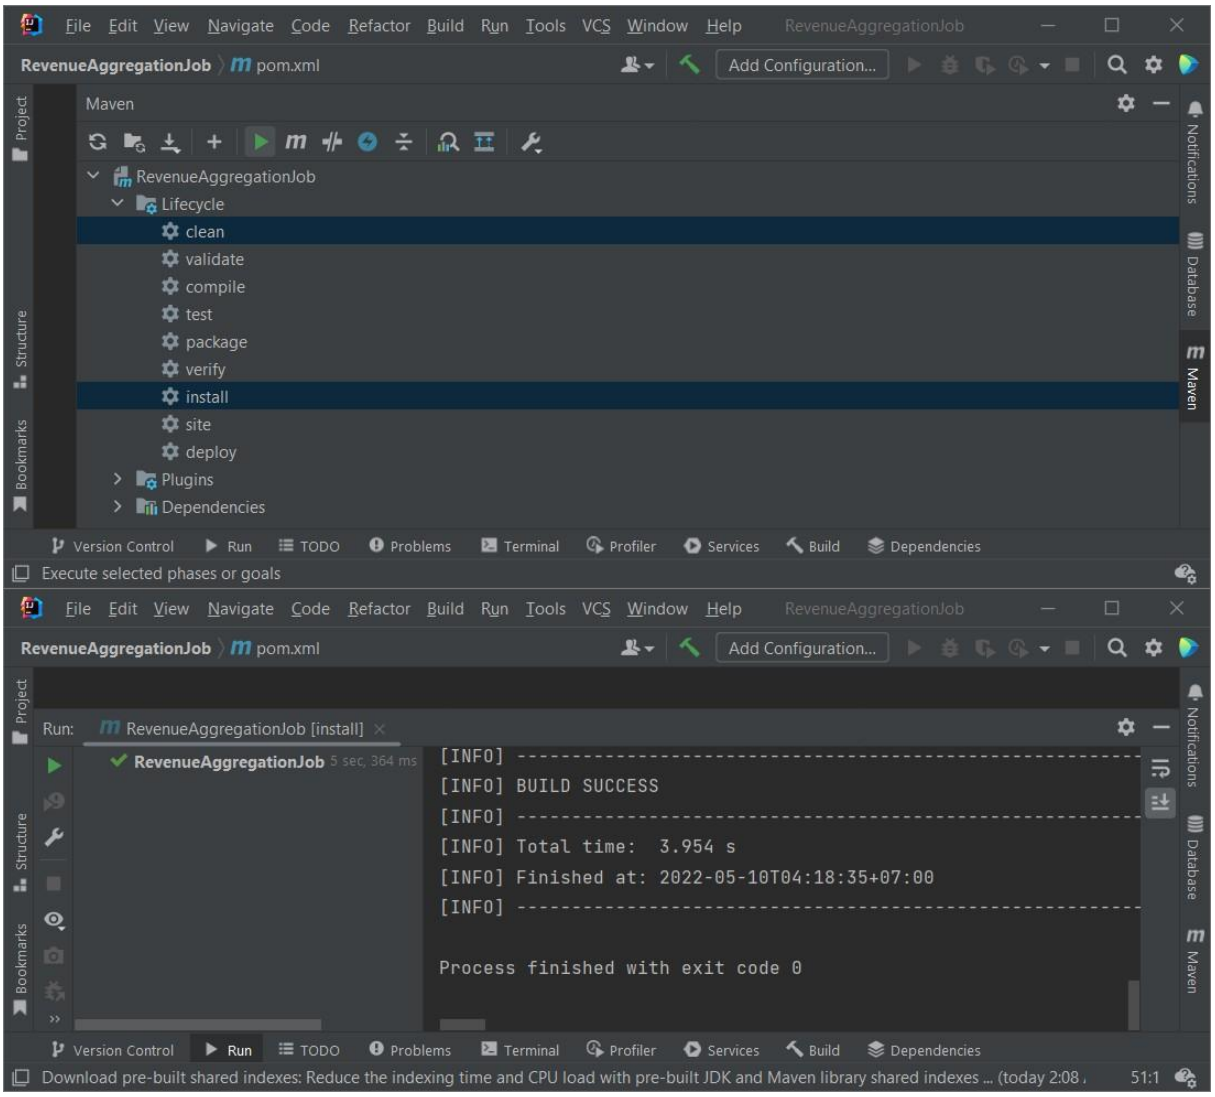

Re-upload the new JAR to the name node.

```
scp RevenueAggregationJob-1.0-SNAPSHOT.jar hadoopuser@192.168.1.14:~/revenue-new.jar
```

```
hadoopuser@192.168.1.14's password:
RevenueAggregationJob-1.0-SNAPSHOT.jar
```

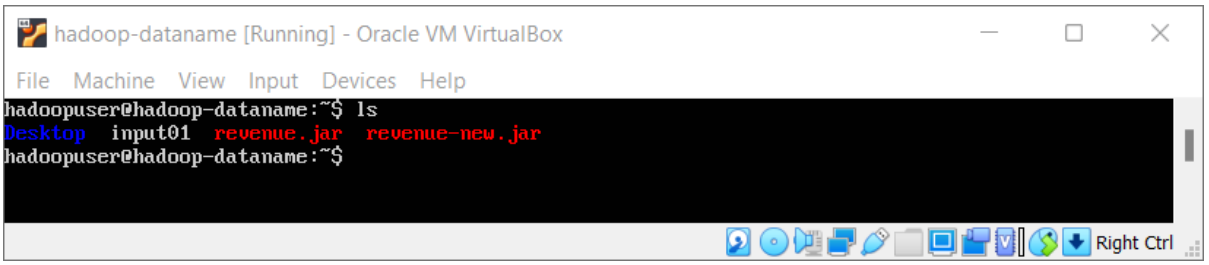

Execution.

```
hadoopuser@hadoop-dataname:~$ hadoop jar revenue-new.jar org.example.App /Output/HasilBaru_
```

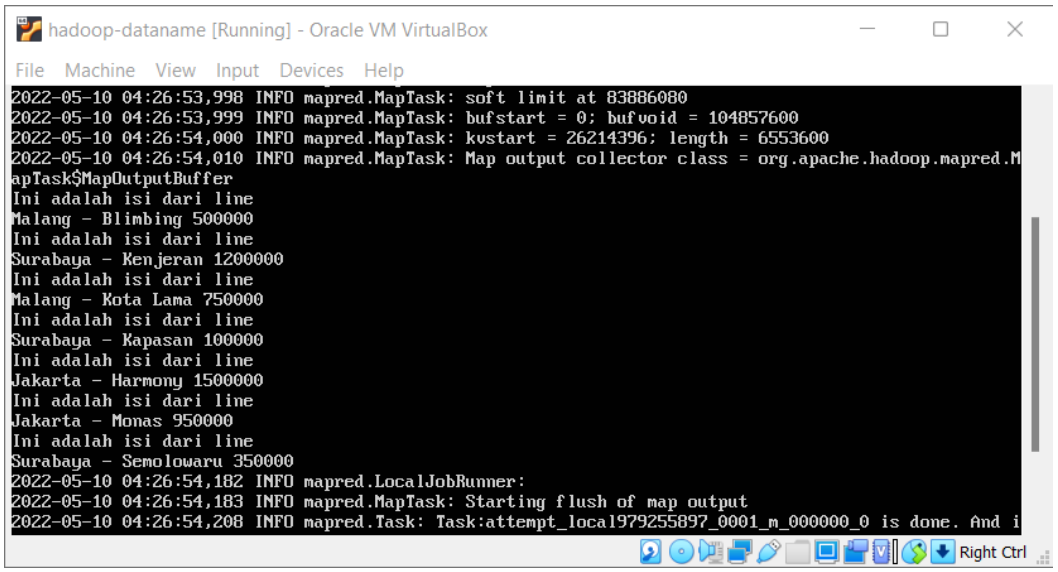

16. Scroll and/or search (Ctrl + F) on your console, is there a bookmark message from the Reducer class?

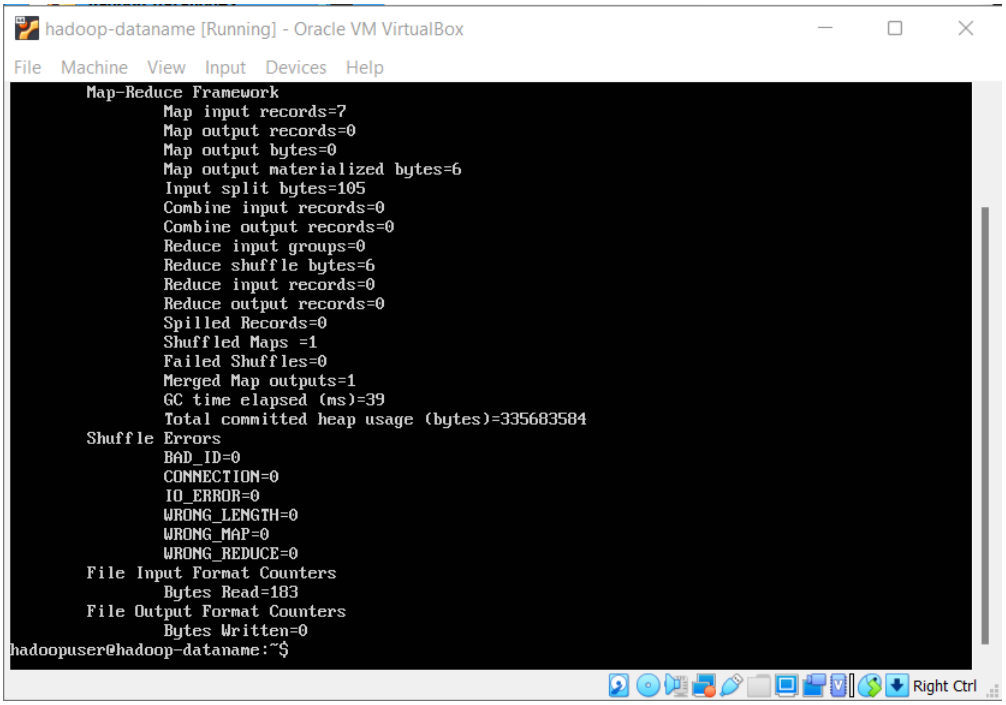

17. Add a new file. The name of the file: "run-job.sh"

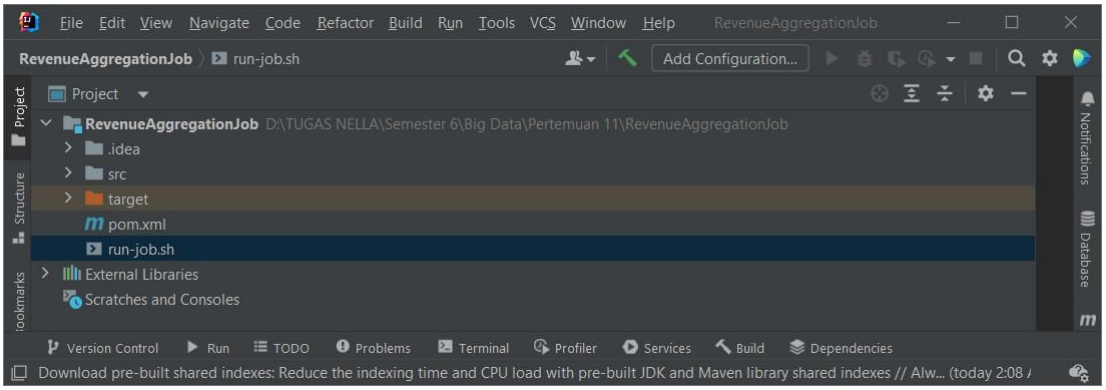

18. Add a script to the file. Adjust the variables in the initial rows of the file to the conditions in the computer and cluster.

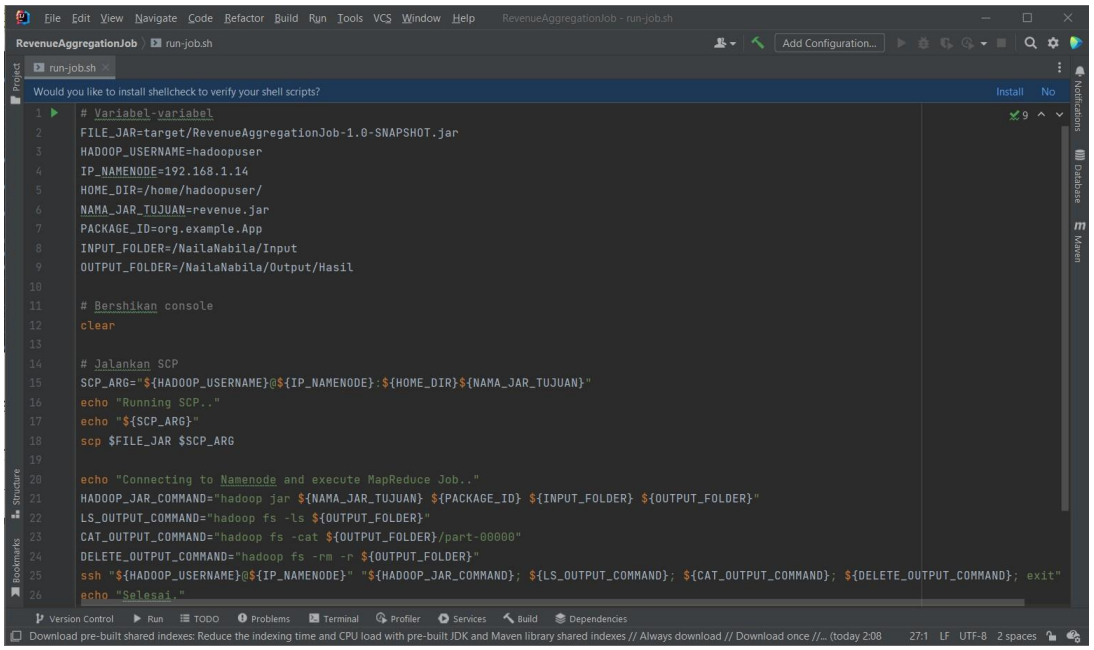

```
# Variabel-variabel
FILE_JAR=target/RevenueAggregationJob-1.0-SNAPSHOT.jar
HADOOP_USERNAME=hadoopuser
IP_NAMENODE=192.168.1.14
HOME_DIR=/home/hadoopuser/
NAMA_JAR_TUJUAN=revenue.jar
PACKAGE_ID=org.example.App
INPUT_FOLDER=/NailaNabila/Input
OUTPUT_FOLDER=/NailaNabila/Output/Hasil

# Bershikan console
clear

# Jalankan SCP
SCP_ARG="${HADOOP_USERNAME}@${IP_NAMENODE}:${HOME_DIR}${NAMA_JAR_TUJUAN}"
echo "Running SCP.."
echo "${SCP_ARG}"
scp $FILE_JAR $SCP_ARG

echo "Connecting to Namenode and execute MapReduce Job.."
HADOOP_JAR_COMMAND="hadoop jar ${NAMA_JAR_TUJUAN} ${PACKAGE_ID} ${INPUT_FOLDER} ${OUTPUT_FOLDER}"
LS_OUTPUT_COMMAND="hadoop fs -ls ${OUTPUT_FOLDER}"
CAT_OUTPUT_COMMAND="hadoop fs -cat ${OUTPUT_FOLDER}/part-00000"
DELETE_OUTPUT_COMMAND="hadoop fs -rm -r ${OUTPUT_FOLDER}"
ssh "${HADOOP_USERNAME}@${IP_NAMENODE}" "${HADOOP_JAR_COMMAND}; ${LS_OUTPUT_COMMAND}; ${CAT_OUTPUT_COMMAND}; ${DELETE_OUTPUT_COMMAND}; exit"
echo "Selesai."
```

19. Execution

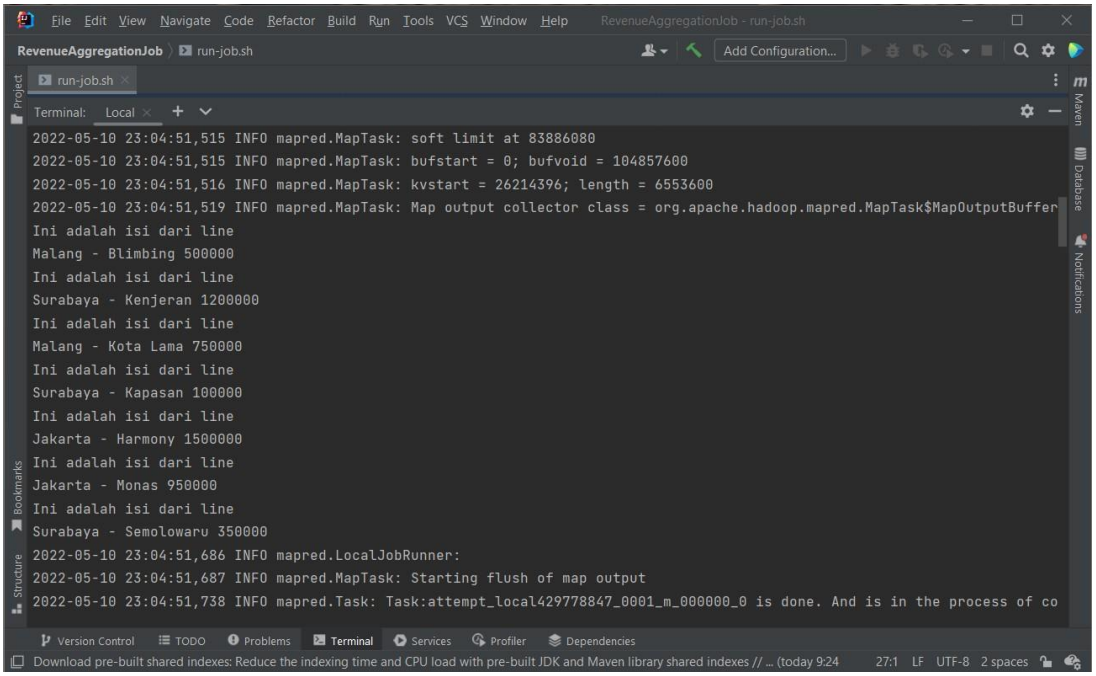

```
2022-05-10 23:04:51,515 INFO mapred.MapTask: soft limit at 83886080
2022-05-10 23:04:51,515 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2022-05-10 23:04:51,516 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2022-05-10 23:04:51,519 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
Ini adalah isi dari line
Malang - Blimbing 500000
Ini adalah isi dari line
Surabaya - Kenjeran 1200000
Ini adalah isi dari line
Malang - Kota Lama 750000
Ini adalah isi dari line
Surabaya - Kapasan 100000
Ini adalah isi dari line
Jakarta - Harmony 1500000
Ini adalah isi dari line
Jakarta - Monas 950000
Ini adalah isi dari line
Surabaya - Semolowaru 350000
2022-05-10 23:04:51,686 INFO mapred.LocalJobRunner:
2022-05-10 23:04:51,687 INFO mapred.MapTask: Starting flush of map output
2022-05-10 23:04:51,738 INFO mapred.Task: Task:attempt_local429778847_0001_m_000000_0 is done. And is in the process of co
```

Part III: Logika Aggregation Process

20. Modify the Class RevenueAggregationMapper.

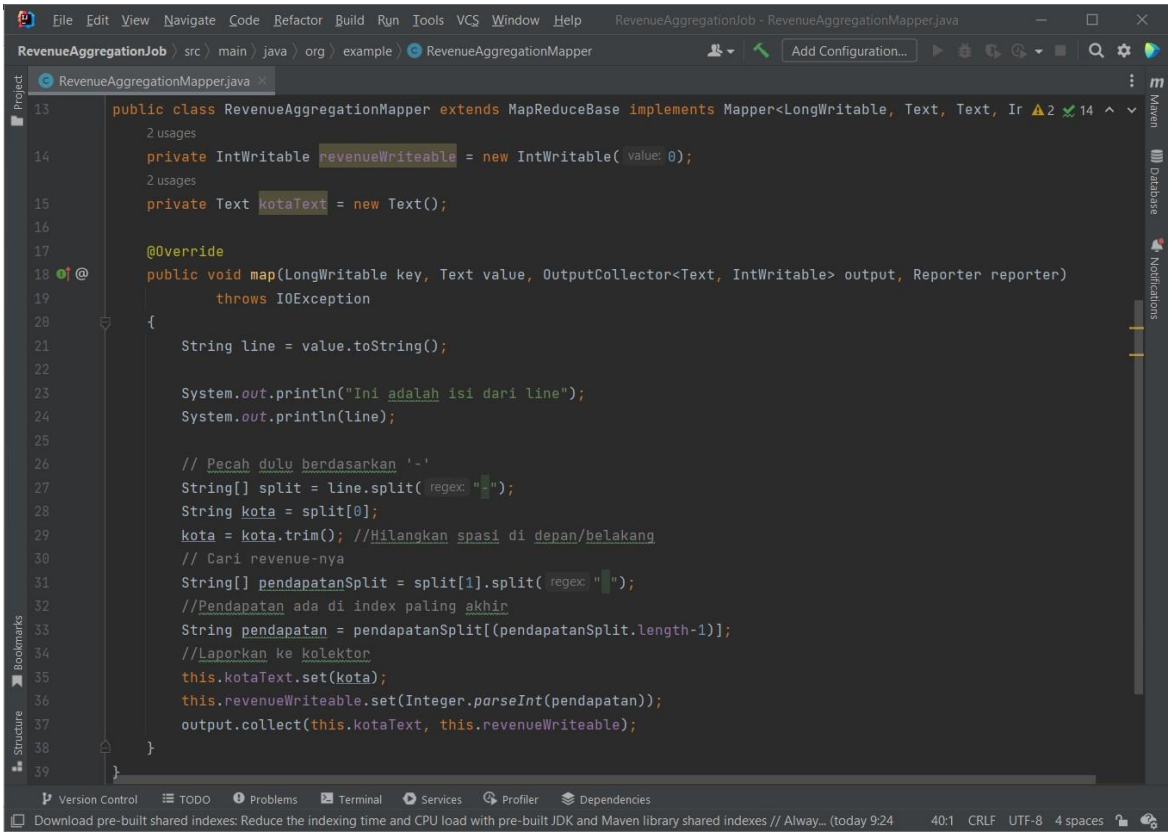

```
public class RevenueAggregationMapper extends MapReduceBase implements Mapper<LongWritable, Text, Text, Ir
    2 usages
    private IntWritable revenueWriteable = new IntWritable( value: 0);
    2 usages
    private Text kotaText = new Text();

    @Override
    public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output, Reporter reporter)
            throws IOException
    {
        String line = value.toString();

        System.out.println("Ini adalah isi dari line");
        System.out.println(line);

        // Pecah dulu berdasarkan '-'
        String[] split = line.split( regex: "-");
        String kota = split[0];
        kota = kota.trim(); //Hilangkan spasi di depan/belakang
        // Cari revenue-nya
        String[] pendapatanSplit = split[1].split( regex: " ");
        //Pendapatan ada di index paling akhir
        String pendapatan = pendapatanSplit[(pendapatanSplit.length-1)];
        //Laporkan ke kolektor
        this.kotaText.set(kota);
        this.revenueWriteable.set(Integer.parseInt(pendapatan));
        output.collect(this.kotaText, this.revenueWriteable);
    }
}
```

21. In the RevenueAggregationReducer class, modify the reduce() method.

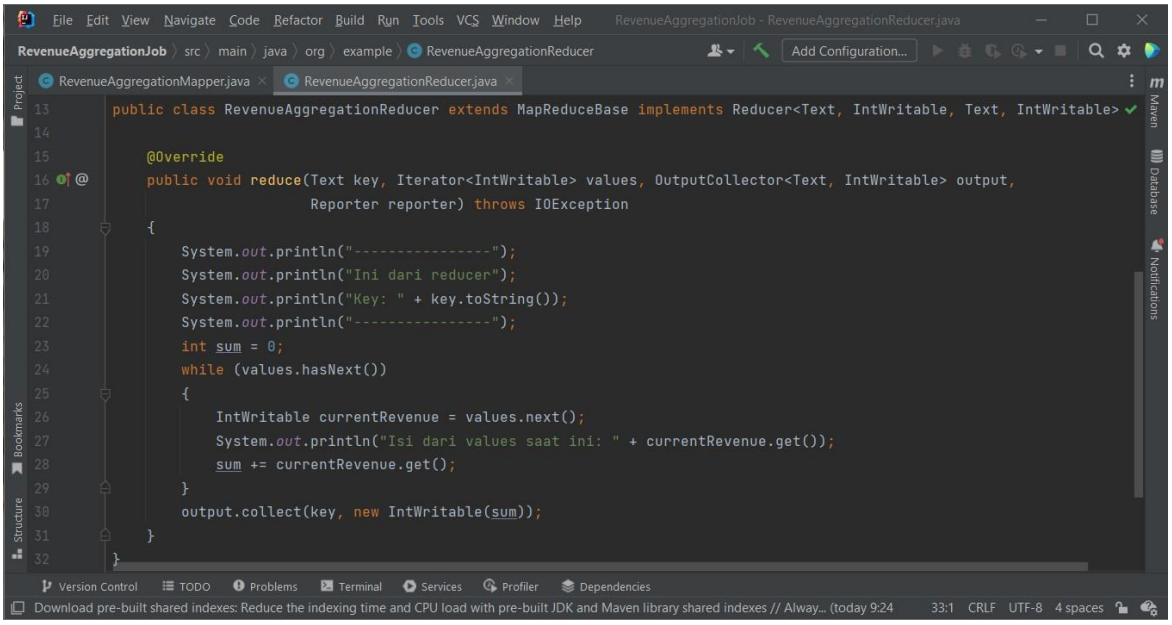

```
public class RevenueAggregationReducer extends MapReduceBase implements Reducer<Text, IntWritable, Text, IntWritable>

    @Override
    public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable> output,
                       Reporter reporter) throws IOException
    {
        System.out.println("----------------");
        System.out.println("Ini dari reducer");
        System.out.println("Key: " + key.toString());
        System.out.println("----------------");
        int sum = 0;
        while (values.hasNext())
        {
            IntWritable currentRevenue = values.next();
            System.out.println("Isi dari values saat ini: " + currentRevenue.get());
            sum += currentRevenue.get();
        }
        output.collect(key, new IntWritable(sum));
    }
}
```

22. Compile, deploy and rerun the MapReduce program and pay attention to the results. Print marker messages from the Reducer class can find them in the console. Which means, the reduce() method in the class, this time called.

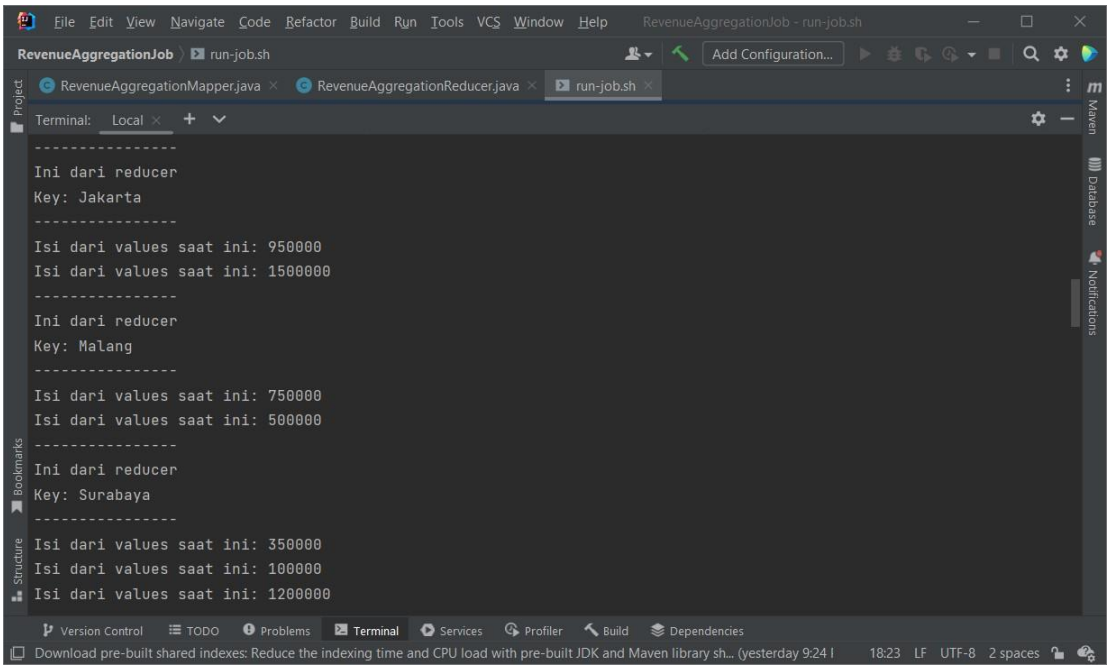

23. Scroll to thebottom of the console, now the results of aggregation should be clearly visible and correct results.

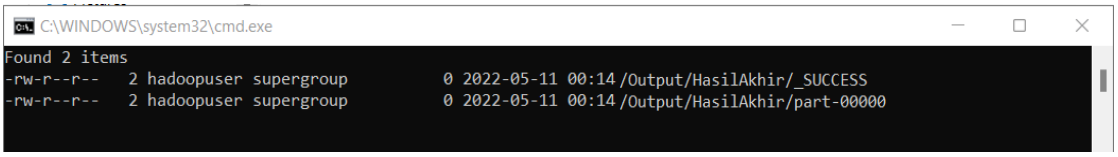

Part I Questions

1. Why is the output file empty/has no content?

   Answer:

   - **Because in the Mapper class and reducer class do not havea command to display the data to be stored, so the process is only run without being stored in the Results directory .**

Question Part II

1. When and how many folders() in the Mapper class are executed?

   Answer:

   - **Method map() is executed when it can be input and the method will be executed as much data from the input.**

2. Why is the reduce() method in the Reducer class never followed?

   Answer:

   - **Because the mapper result data is not sent into the Reducer class so the reduce() method does not have the data to be executed.**

Question Part III

1. When and how much is the reduce() method run?

   Answer:

   - **The method will be run when it can be input from the Mapper class and run 3 times because there are 3 keys, namely Malang, Surabaya, and Jakarta.**

2. Why can the console appear the contents of part-0000 files?

   Answer:

   - **Because the reducer result was successfully stored on the part-0000 file.**

Assignment

Try to replace its aggregation operations to be average per city instead of counting the totals!

Task Result:

- Step 1: Modification of the RevenueAggregationReducer class

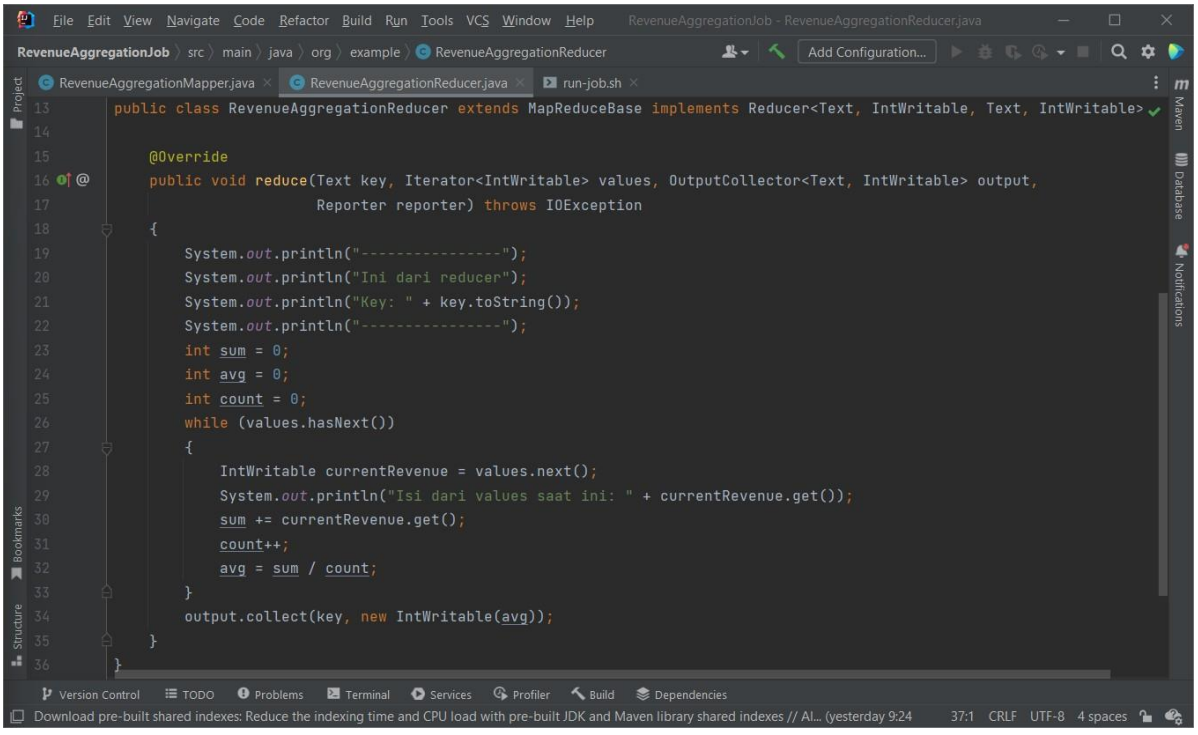

- Step 2: Compile and Build Project

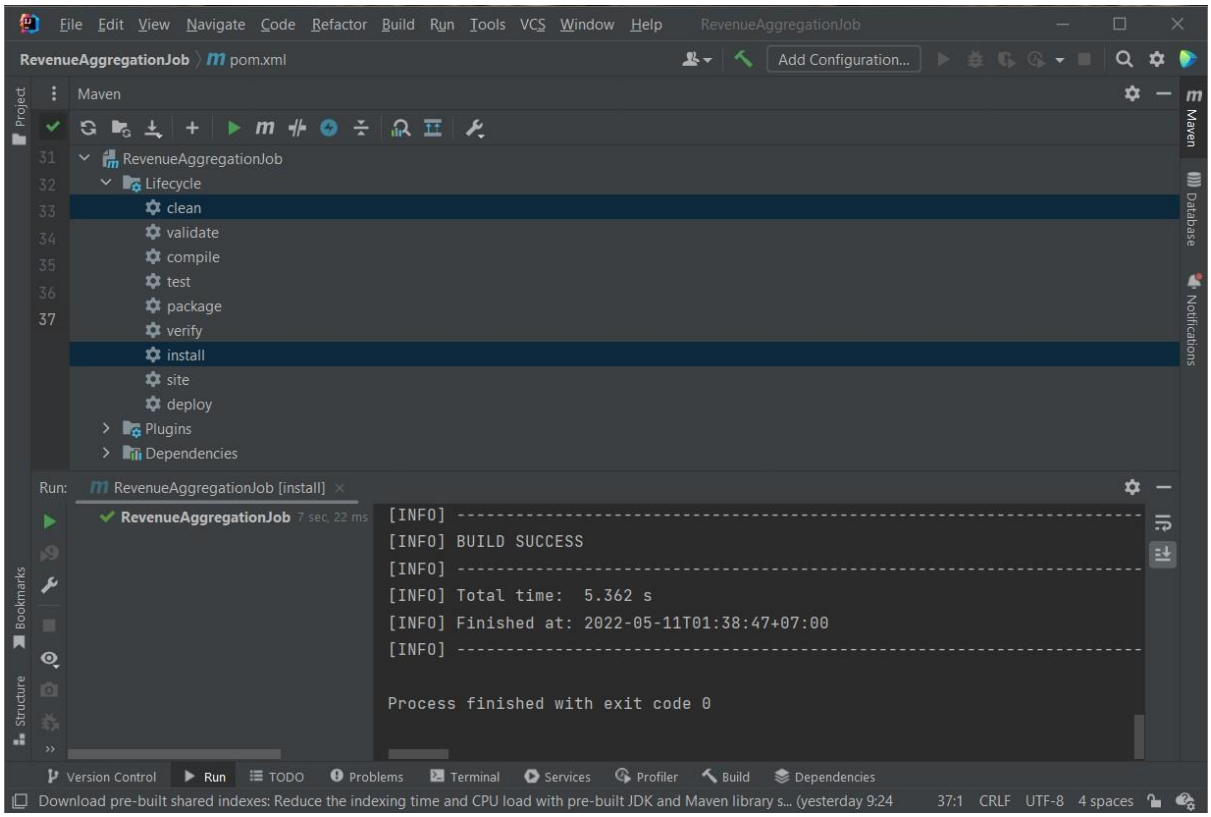

- Startl

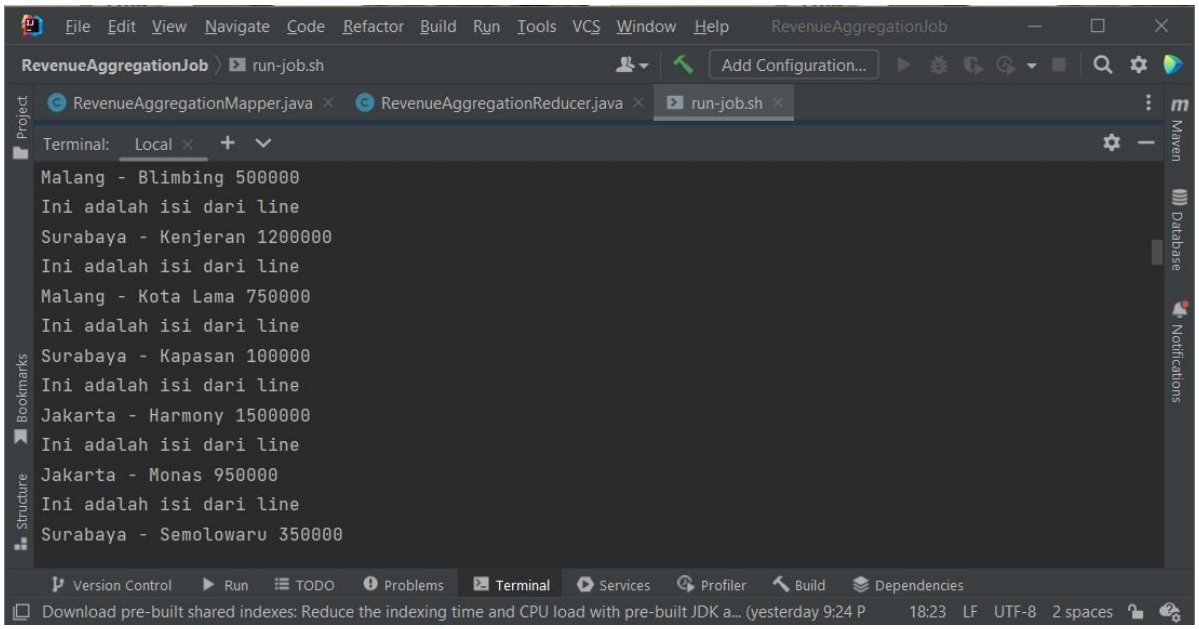

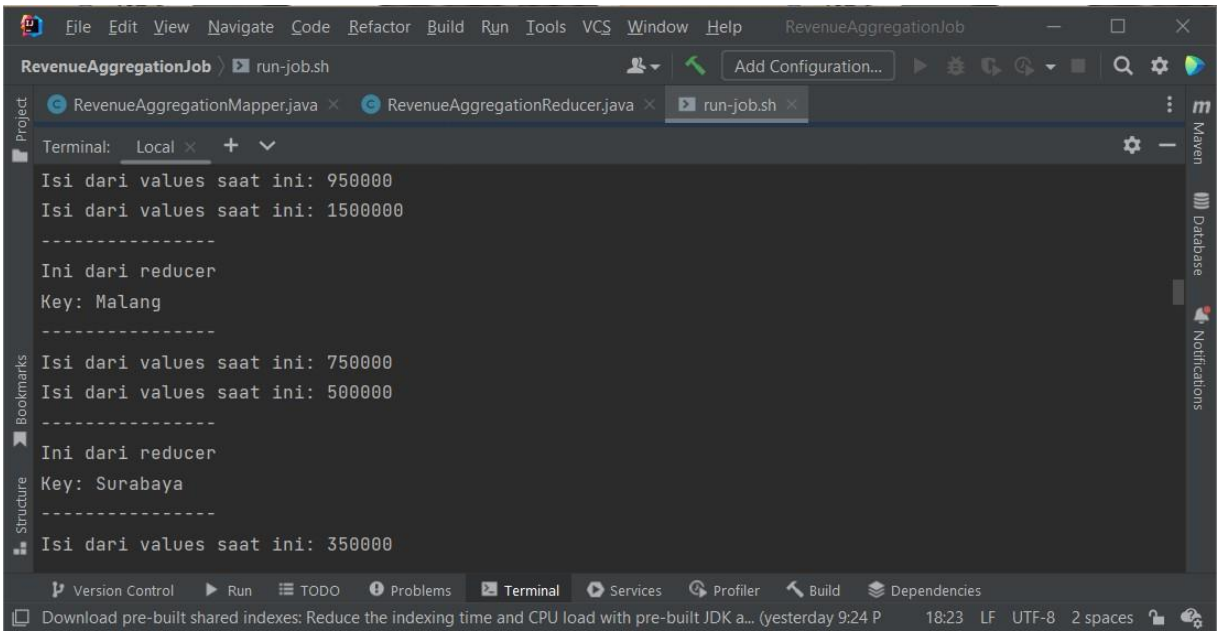
File Edit View Navigate Code Refactor Build Run Tools VCS Window Help
RevenueAggregationJob
RevenueAggregationJob 〉 run-job.sh
Add Configuration...
RevenueAggregationMapper.java
RevenueAggregationReducer.java
run-job.sh
Terminal: Local
Isi dari values saat ini: 950000
Isi dari values saat ini: 1500000

Ini dari reducer
Key: Malang

Isi dari values saat ini: 750000
Isi dari values saat ini: 500000

Ini dari reducer
Key: Surabaya

Isi dari values saat ini: 350000
Version Control
Run
TODO
Problems
Terminal
Services
Profiler
Build
Dependencies
Download pre-built shared indexes: Reduce the indexing time and CPU load with pre-built JDK a... (yesterday 9:24 P
18:23
LF
UTF-8
2 spaces

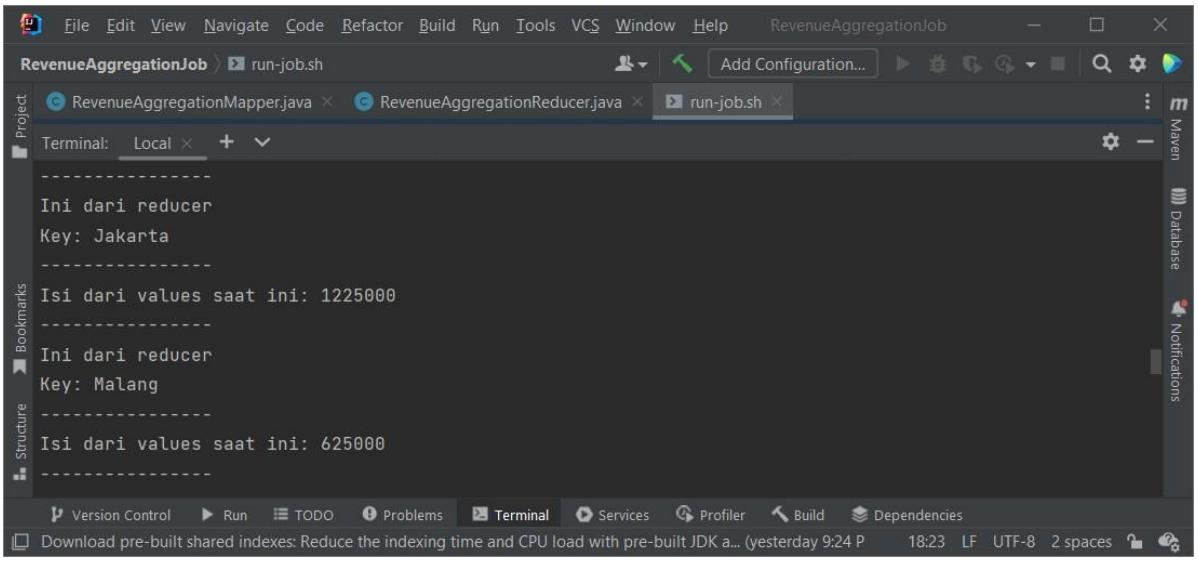
File Edit View Navigate Code Refactor Build Run Tools VCS Window Help
RevenueAggregationJob
RevenueAggregationJob 〉 run-job.sh
Add Configuration...
RevenueAggregationMapper.java
RevenueAggregationReducer.java
run-job.sh
Terminal: Local

Ini dari reducer
Key: Jakarta

Isi dari values saat ini: 1225000

Ini dari reducer
Key: Malang

Isi dari values saat ini: 625000

Version Control
Run
TODO
Problems
Terminal
Services
Profiler
Build
Dependencies
Download pre-built shared indexes: Reduce the indexing time and CPU load with pre-built JDK a... (yesterday 9:24 P
18:23
LF
UTF-8
2 spaces